



STEAM ANALYTICS: REAL-TIME TELEMETRY & CRISIS DETECTION

Group 22

20225473 - Nguyễn Trọng Phương Bách

20225494 - Vũ Minh Hiếu

20225501 - Bùi Nguyên Khải

20225523 - Trần Ngọc Quang

20225533 - Nguyễn Việt Tiến

ONE LOVE. ONE FUTURE.

The Problem Context

The Scale

- Steam Platform: 50,000+ games, 120M+ monthly active users.
- Data Velocity: Thousands of reviews and player status updates every second.

The Big Data Challenge (4 Vs)

- Volume: Massive datasets of historical and incoming reviews.
- Velocity: Real-time demands; waiting for daily batches is too slow.
- Variety: structured, semi-structured, time-series
- Verification: Need to handle strict API rate limits.

Core Issue

- Traditional metrics are reactive.
- Developers often miss "Review Bombing" attacks (coordinated negative campaigns) until the damage is done.



Project Goals & Innovation

Objective

- Build a production-grade, containerized real-time analytics pipeline.
- Goal: Reduce insight latency from 24 hours to <1 minute.

Key Innovation: Review Bomb Detection

- Algorithm: Real-time flagging of "Panic Events".
- Trigger Formula:


$$\frac{R_{\text{negative}}}{R_{\text{total}}} > 0.8 \quad \text{AND} \quad R_{\text{total}} > 10$$

User Scenario

- **Before:** Team wakes up to 1,000 negative reviews. Damage done.
- **After:** Dashboard Alert triggers in **30 seconds**. Immediate rollback.

System Architecture Overview

Real-Time Steam Analytics Pipeline (Kappa Architecture)



Technology Stack

Platform: Kubernetes (Docker Desktop).

Management: Stackable Data Platform

-Core Infrastructure Components:

- **Apache Zookeeper:** The "Brain" of the cluster.
 - Role: Distributed coordination, Leader Election for Kafka, and High Availability (HA) for HDFS NameNodes.
- **Stackable Operators:**
 - Role: Automates the complex lifecycle (deployment, scaling, upgrades) of Zookeeper, Kafka, and HDFS on Kubernetes.
- **Streaming Engine:** Spark 3.5.6 + Kafka 3.9.1.
- **Visualization:** Grafana + Prometheus + MongoDB Express.



Data Ingestion

- Data producer:
 - **Reviews & Info:** Python-based Kubernetes CronJob (producer_reviews.py, producer_charts.py) that scrapes app details and user reviews.
 - **Player Counts:** Python-based Kubernetes CronJob (producer_players.py) running every 5 minutes to fetch real-time player statistics.
- Kafka Cluster Structure:
 - **Deployment:** Hosted on Kubernetes using the Stackable Operator.
 - **Configuration:**
 - **Single Broker** setup optimized for development/testing (1 replica).
 - **Retention Policy:** 7 days (168 hours) to support Kappa architecture replayability.
 - **Partitions:** 3 Partitions per topic.
 - **Resources:** Limited to 1 CPU and 2Gi Memory per broker.
 - **Zookeeper:** Integrated for cluster coordination.
- **Transport Encryption:** All traffic between producers and the Kafka cluster is encrypted using **SSL/TLS**.

Data Ingestion

- Kafka message structure:

- Player count:

```
player_record = {
    "appid": int(appid),
    "player_count": p_count,
    "timestamp": datetime.utcnow().isoformat()
}
```

- Game info:

```
return {
    "appid": int(appid),
    "name": item.get("name") or data.get("name"),
    "primary_genre": genre_tag,
    "type": data.get("type"),
    "release_date": data.get("release_date", {}).get("date"),
    "is_free": data.get("is_free"),
    "short_description": clean_html(data.get("short_description")),
    "developers": data.get("developers") or [],
    "publishers": data.get("publishers") or [],
    "genres": [g.get("description") for g in (data.get("genres") or [])],
    "price_overview": data.get("price_overview") or {},
    "categories": [c.get("description") for c in (data.get("categories") or [])],
    "metacritic": data.get("metacritic", {}).get("score"),
    "recommendations": data.get("recommendations", {}).get("total"),
    "achievements_count": data.get("achievements", {}).get("total", 0),
    "timestamp_scraped": datetime.utcnow().isoformat()
}
```

Data Ingestion

- Kafka message structure:
 - Game review:

```
yield {
    "app_id": app_id,
    "review_id": r.get("recommendationid"),
    "author_steamid": r.get("author", {}).get("steamid"),
    "playtime_at_review": r.get("author", {}).get("playtime_at_review"),
    "playtime_forever": r.get("author", {}).get("playtime_forever"),
    "language": r.get("language"),
    "voted_up": r.get("voted_up"),
    "votes_up": r.get("votes_up"),
    "weighted_vote_score": r.get("weighted_vote_score"),
    "timestamp_created": r.get("timestamp_created"),
    "review_text": clean_html(r.get("review")),
    "scraped_at": datetime.utcnow().isoformat()
}
```

Stream Processing Logic

- Charts App
 - **Genre Explosion:** Unnests arrays (e.g., [Action, RPG]) via `explode` for accurate category counts.
- Reviews App
 - **Windowing:** Uses **1-hour tumbling windows** with a **10-minute watermark** to manage late data .
 - **Metrics:** Aggregates total_reviews and avg quality to track sentiment .
- Players App
 - **Granularity:** Uses **10-minute windows** to capture rapid activity fluctuations.
- Resilience Features
 - **Schema Validation:** Strict `StructType` parses invalid JSON as `null` to prevent crashes.
 - **Watermarking:** Discards records exceeding the 5-minute threshold to prevent memory growth.
 - **Checkpointing:** Stores query state in dedicated HDFS directories (e.g., `/checkpoints/reviews_cold`) to support restartability.

Storage Strategy

- **The Dual-Write Approach:**
 - Cold Storage – Bronze Layer (Raw data): Apache HDFS
 - Purpose: Archival, long-term retention
 - Format: Parquet (Columnar storage).
 - Benefit: 10:1 Compression ratio compared to raw JSON.
 - Path: `/archive/reviews`, `/archive/charts`, `/archive/players`.
 - Use Case: DL/ML models training or debugging pipeline logic.
 - Hot Storage – Gold Layer (Aggregated data): MongoDB
 - Purpose: Business-level aggregates ready for consumption (operational analytics, dashboards, web app API)
 - Format: Optimized Documents with Arrays/Counters.
 - Optimization: TTL Indexes (30 days) to prevent storage bloat and maintain query speed.
 - Use case: Powering the Real-Time Grafana Dashboard.
 - Note: We skip a persistent "Silver" layer to minimize latency (Stream-to-Gold).

HDFS Schema (Parquet)

- **Reviews**
 - app_id (String)
 - review_id (String)
 - recommended (Boolean)
 - votes_up (Integer)
 - weighted_vote_score (Float)
 - timestamp (Timestamp) :: Event Time
- **Charts**
 - appid (Integer)
 - name (String)
 - type (String) :: e.g., "game", "dlc"
 - genres (Array[String]) :: e.g. ["Action", "FPS"]
 - timestamp (Timestamp) :: Scrape Time
- **Players**
 - appid (Integer)
 - player_count (Integer)
 - timestamp (Timestamp) :: Observation Time

Visualization

- **Spark UI:** A comprehensive interface to monitor and analyze the execution of Spark applications and offers insights into the status and resource consumption of Spark cluster.

The screenshot shows the Apache Spark 3.5.6 UI interface. At the top, there is a navigation bar with tabs: Jobs (which is selected), Stages, Storage, Environment, Executors, SQL / DataFrame, and Structured Streaming. The main content area is titled "Spark Jobs" with a link to the documentation. It displays the following information:

- User: stackable
- Total Uptime: 3.2 min
- Scheduling Mode: FIFO
- Completed Jobs: 3

Below this, there are two expandable sections: "Event Timeline" and "Completed Jobs (3)". The "Completed Jobs" section is expanded, showing three completed jobs with their IDs and descriptions:

Job Id (Job Group) ▾	Description
2 (5d39dee5-62c1-427c-9087-04912534a4e8)	id = ed34313f-920c-4e46-b1c7-b1eb632ec05a runId = 5d39dee5-62c1-427c-9087-04912534a4e8 batch = 0 start at NativeMethodAccessorImpl.java:0
1 (92839a4c-54af-4471-ae6a-36ea2493122c)	id = 224e5950-7ef2-4df9-9ffc-11ceedbf971d runId = 92839a4c-54af-4471-ae6a-36ea2493122c batch = 0 start at NativeMethodAccessorImpl.java:0
0 (6a6449c5-020e-46ad-a552-2ab55e347f21)	id = d169c8a5-5972-4006-a54a-9599a1cb9904 runId = 6a6449c5-020e-46ad-a552-2ab55e347f21 batch = 0 start at NativeMethodAccessorImpl.java:0

Visualization

- **Mongo Express:** A web-based admin interface for browsing MongoDB collections and running ad-hoc queries.

The screenshot shows the Mongo Express interface for the 'bigdata' database. At the top, there's a header with a green circle icon, the text 'Mongo Express Database: bigdata', and a dropdown menu. Below the header, the title 'Viewing Database: bigdata' is displayed. Underneath, the 'Collections' section lists four entries:

Collection Name	+ Create collection
review_bomb_alerts	
steam_charts	
steam_players	
steam_reviews	

Each collection entry includes a 'View' button (green), 'Export' (orange), '[JSON]' (orange), and 'Import' (orange) button. To the right of the collection names are red 'Del' buttons. Below this section is a 'Database Stats' table:


Database Stats	
Collections (incl. system.namespaces)	4
Data Size	264 KB
Storage Size	188 KB
Avg Obj Size #	116 Bytes
Objects #	2273
Indexes #	4
Index Size	156 KB

- **Grafana:** An open-source analytics and monitoring platform that let users query, visualize, and alert on metrics from various data sources.
- Usage:
 - Monitor system metrics of databases (MongoDB and Prometheus).
 - Display key features of collected Steam data.




Visualization

- Monitor system metrics of databases:
 - Prometheus collects and stores time-series metrics (such as CPU/memory usage, request counts) from systems and applications.
 - Information is shown in Grafana's Drilldown section.



Visualization

- Display key features of collected Steam data:
 - **Review Bomb alert:** Green/red status indicator of abnormal surge in negative reviews.
 - **Top games:** Show top games by number of players across Steam.
 - **Performance check:** Real-time concurrent player counts.
... and more.



Lesson Learned

1. The "TLS Nightmare" (Security Mismatch)

- Context: Stackable enforces TLS, but Spark Java clients expect JKS/PKCS12.
- Fix: Implemented Kubernetes Init Containers to run keytool conversion before the main application starts.

2. The "Small Files" Problem (HDFS)

- Context: Real-time streaming generates thousands of tiny KB-sized files, threatening NameNode stability.
- Fix: Increased HDFS write trigger to 1 minute (Micro-batching) to reduce file creation frequency by 60x compared to continuous processing.

3. Running a Data Center on a Laptop (Resources)

- Context: Running Zookeeper, Kafka, HDFS, Spark, and Mongo on Docker Desktop (8GB RAM limit).
- Fix: Aggressive tuning:
 - Reduced Zookeeper/NameNodes to 512MB.
 - Allocated 2GB to Kafka (for buffering).
 - Running single replicas instead of full HA for dev.

Business Impact & Conclusion

- **Speed:** Insight generation reduced from 24+ hours to 30 seconds.
- **Efficiency:** Replaced manual analyst monitoring with automated alerts.
- **Scalability:** Zero-downtime scaling via Kubernetes replicas.
- **Path to Production Scale**
 - Storage: Upgrade HDFS to Delta Lake for ACID transactions and automatic small-file compaction.
 - Infrastructure: Migrate to Managed K8s (EKS/AKS) to scale workers dynamically.
 - Governance: Add Confluent Schema Registry to enforce strict API contracts.
 - Dashboard: Hype-to-Disappointment Index (predict refund waves during a launch)

Conclusion: Kappa Architecture is a practical, robust, and efficient solution for handling the velocity and volume of modern gaming telemetry.



HUST

THANK YOU !